

Implement a Basic Driving Agent

To begin, your only task is to get the **smartcab** to move around in the environment. At this point, you will not be concerned with any sort of optimal driving policy. Note that the driving agent is given the following information at each intersection:

- The next waypoint location relative to its current location and heading.
- The state of the traffic light at the intersection and the presence of oncoming vehicles from other directions.
- The current time left from the allotted deadline.

To complete this task, simply have your driving agent choose a random action from the set of possible actions (**None**, **'forward'**, **'left'**, **'right'**) at each intersection, disregarding the input information above. Set the simulation deadline enforcement, **enforce_deadline** to **False** and observe how it performs.

QUESTION: *Observe what you see with the agent's behavior as it takes random actions. Does the **smartcab** eventually make it to the destination? Are there any other interesting observations to note?*

ANSWER: *The agent just takes random actions. The goal is reached; however, this happens in an arbitrary amount of time. The agent can be in the previous waypoint but just take a different direction than the one of the goal. Deadline goes to 0 and then to negative values, and reduction in the deadline value does not stop until goal is reached and the board is reset.*

Inform the Driving Agent

Now that your driving agent is capable of moving around in the environment, your next task is to identify a set of states that are appropriate for modeling the **smartcab** and environment. The main source of state variables are the current inputs at the intersection, but not all may require representation. You may choose to explicitly define states, or use some combination of inputs as an implicit state. At each time step, process the inputs and update the agent's current state using the **self.state** variable. Continue with the simulation deadline enforcement **enforce_deadline** being set to

False, and observe how your driving agent now reports the change in state as the simulation progresses.

QUESTION: What states have you identified that are appropriate for modeling the **smartcab** and environment? Why do you believe each of these states to be appropriate for this problem?

OPTIONAL: How many states in total exist for the **smartcab** in this environment? Does this number seem reasonable given that the goal of Q-Learning is to learn and make informed decisions about each state? Why or why not?

ANSWER: The states selected are {inputs, self.next_waypoint}. I chose them because I believe this way we can represent all the information we have with regards to the current state. To my understanding, since we do not have our global location, the taxi determines its state by what the traffic light is, where the planner says should be the next waypoint, and what the intended move is, for vehicles that are oncoming, or coming from the left or the right side of the intersection.

Thus, if I have not misunderstood, the taxi does not take into consideration at which intersection it is (the taxi does not know), and treats all points where 'light', 'oncoming', 'right', 'left', and 'waypoint' are the same, as one. For all these points, if there is no random selection chosen, the same action will be chosen (while, of course, $Q(S,a)$ remains the same).

'light' can take 2 values, 'oncoming', 'right', and 'left' can take 4 each, and 'waypoint' can take 3. Thus, the possible states in total are $2*4*4*4*3 = 384$

The deadline also provides information, however has many values, and thus it would increase the number of possible states significantly. (For 100 different values we would then have 3840 states). So I chose not to include it in the state.

A number of 384 states, in my opinion is reasonable. In the specific application that we work on, since the Q table is not reset for every new destination, we will probably have evaluated all state/action pairs before the program exits.

Implement a Q-Learning Driving Agent

With your driving agent being capable of interpreting the input information and having a mapping of environmental states, your next task is to implement the Q-Learning algorithm for your driving agent to choose the *best* action at each time step, based on the Q-values for the current state and action. Each action taken by the **smartcab** will produce a reward which depends on the state of the environment. The Q-Learning driving agent will need to consider these rewards when updating the Q-values. Once implemented, set the simulation deadline enforcement `enforce_deadline` to `True`. Run the simulation and observe how the **smartcab** moves about the environment in each trial.

The formulas for updating Q-values can be found in [this](#) video.

QUESTION: *What changes do you notice in the agent's behavior when compared to the basic driving agent when random actions were always taken? Why is this behavior occurring?*

ANSWER: *I believe that the agent now does not explore as much as before. Some actions tend to repeat themselves when the agent finds itself in the same intersections, and sometimes the agent tends to seclude itself in a smaller area of the grid, while before it seemed to traverse more waypoints.*

I believe that this occurs because the agent was given some reward for some states that is high, so being in the same state, the agent tends to repeat the action that yielded this high reward.

Improve the Q-Learning Driving Agent

Your final task for this project is to enhance your driving agent so that, after sufficient training, the **smartcab** is able to reach the destination within the allotted time safely and efficiently. Parameters in the Q-Learning algorithm, such as the learning rate (`alpha`), the discount factor (`gamma`) and the exploration rate (`epsilon`) all contribute to the driving agent's ability to learn the best action for each state. To improve on the success of your **smartcab**:

- Set the number of trials, `n_trials`, in the simulation to 100.
- Run the simulation with the deadline enforcement

`enforce_deadline` set to `True` (you will need to reduce the update delay `update_delay` and set the `display` to `False`).

- Observe the driving agent's learning and **smartcab's** success rate, particularly during the later trials.
- Adjust one or several of the above parameters and iterate this process.

This task is complete once you have arrived at what you determine is the best combination of parameters required for your driving agent to learn successfully.

QUESTION: Report the different values for the parameters tuned in your basic implementation of Q-Learning. For which set of parameters does the agent perform best? How well does the final driving agent perform?

ANSWER: I run the program 100 times (with 100 simulations each), for various alpha, gamma and epsilon combinations. I modified the code to do this automatically, export the results to an excel file, and print out the best result.

The range of alpha and gamma was (0 to 1 – 1 included) with steps of 0.1, and the range of epsilon was 0.0 to 0.15 with a step of 0.05. The best result was reported for alpha = 0.6, gamma = 0.6, and epsilon = 0.1.

The success percentage (car finding the goal) however, was extremely low, near 1.5%, which suggests I am doing something very wrong...

QUESTION: Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties? How would you describe an optimal policy for this problem?

ANSWER: Since the success rate of the taxi seems to be very low, I would say that the agent did not get close to finding a policy